



FURTHER ADVENTURES IN UNCERTAINTY QUANTIFICATION: CONFORMAL PREDICTIONS

Viviana Acquaviva
CUNY/Columbia University

Holy Grail of Uncertainty Quantification:

Generate predictions that are calibrated *for every example*

e.g., guarantee that true value
is within quoted “90%” interval
90% of the time
for any input value

Conformal predictions

Goal: starting from predictions for any model, generate **calibrated** sets (classification) or intervals (regression).

- **distribution-free**: the only assumption is that data points are **exchangeable (no time series)**.
- **model-agnostic**: conformal predictions can be applied to **any** predictive model.
- **marginal coverage guarantee**: the resulting prediction sets (intervals) come with **guarantees of covering the true outcome with at least desired probability** on a set with the same statistical properties

Process

Divide data in train / calibration / validation (test) set

Build model on **train** set; generate predictions on **calibration** set

Use calibration set to build **non-conformity** (or **conformal**) scores

Can be defined in many ways (important: **good model** \Leftrightarrow **low scores**)

In practice, they measure how "off" model predictions are by looking at some statistical property of the scores' distribution

Use non-conformity scores to **generate new prediction sets/intervals** C that satisfy coverage guarantee for any chosen probability $1 - \alpha$:

$$P(Y_{val}) \in C(X_{val}) \geq (1 - \alpha)$$

Classification example

- Let's say the true class of object i is y_i
- Start from any model that predicts **probabilities** $f(x_i)$ for each class
(for example, $f(x_i)$ could be 0.2, 0.7, 0.1 in a three-class model)
- On calibration set, compute scores $s_i = 1 - f(x_i)[y_i]$, i.e.
 $1 - (\text{probability assigned to correct class})$
 s_i are a simple example of **conformal scores** (low = good)
- Sort from certain (low s_i) to uncertain
- Select **probability threshold of interest** q (e.g. 80%)
- Find “ q ” quantile of s_i , multiply by finite sample size correction ($\hat{q} = q * (n+1)/n$)
- Prediction output is the set of all classes that had a probability score $> 1 - \hat{q}$
- Marginal coverage: on statistically equivalent (i.i.d.) data, **sets contain the true class 80% of the time**

	$p(y_1)$	$p(y_2)$	$p(y_3)$
#1	0.2	0.7	0.1
#2	0.4	0.4	0.2
#3	0.9	.05	.05
#4	0.1	0.1	0.8
#5	0.2	0.6	0.2

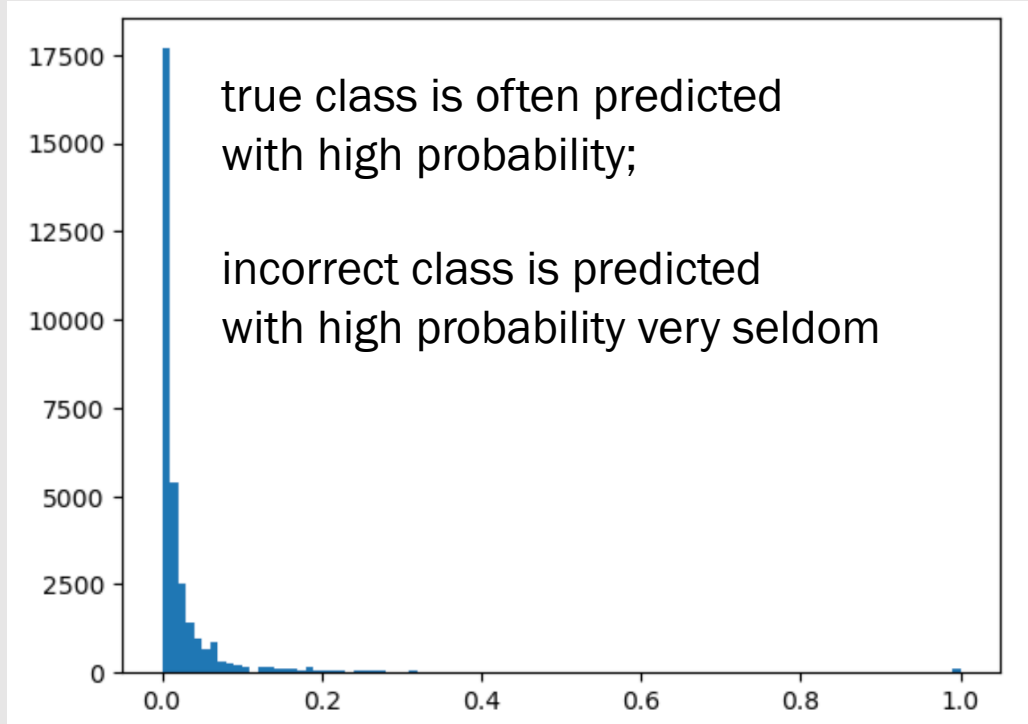
s_i	s_i (sorted)	$q = 0.6$ $\hat{q} = 0.72$
0.3	0.1	
0.6	0.2	
0.1	0.3	
0.2	0.6	
0.8	0.8	

pred sets
2
1,2
1
3
2

How can we get a coverage guarantee
no matter the quality of the estimator?

Good Classifier

conformal scores

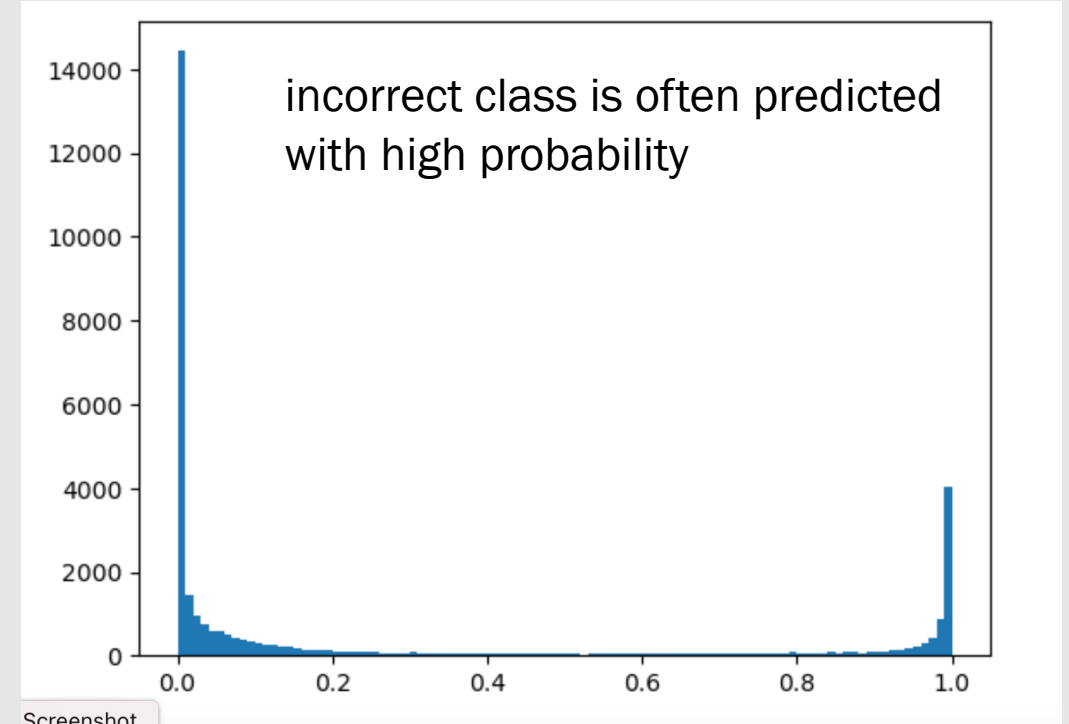


\hat{q} (value corresponding to say 80% quantile of conformal scores) is **low**

only classes predicted with **high probability** ($1 - \hat{q}$) are included in prediction sets

Bad Classifier

conformal scores



\hat{q} (value corresponding to say 80% quantile of conformal scores) is **high**

even classes predicted with **low probability** ($1 - \hat{q}$) are included in prediction sets

Devil is in the detail

- Marginal coverage guarantee **doesn't make** single predictions **particularly useful**

Marginal means “on average”: given a new set of examples that is statistically equivalent to the calibration set, the predictions sets contain the true value “x”% of the time

- What we really want is **conditional coverage**: conditional on any split of the data, i.e. for every example

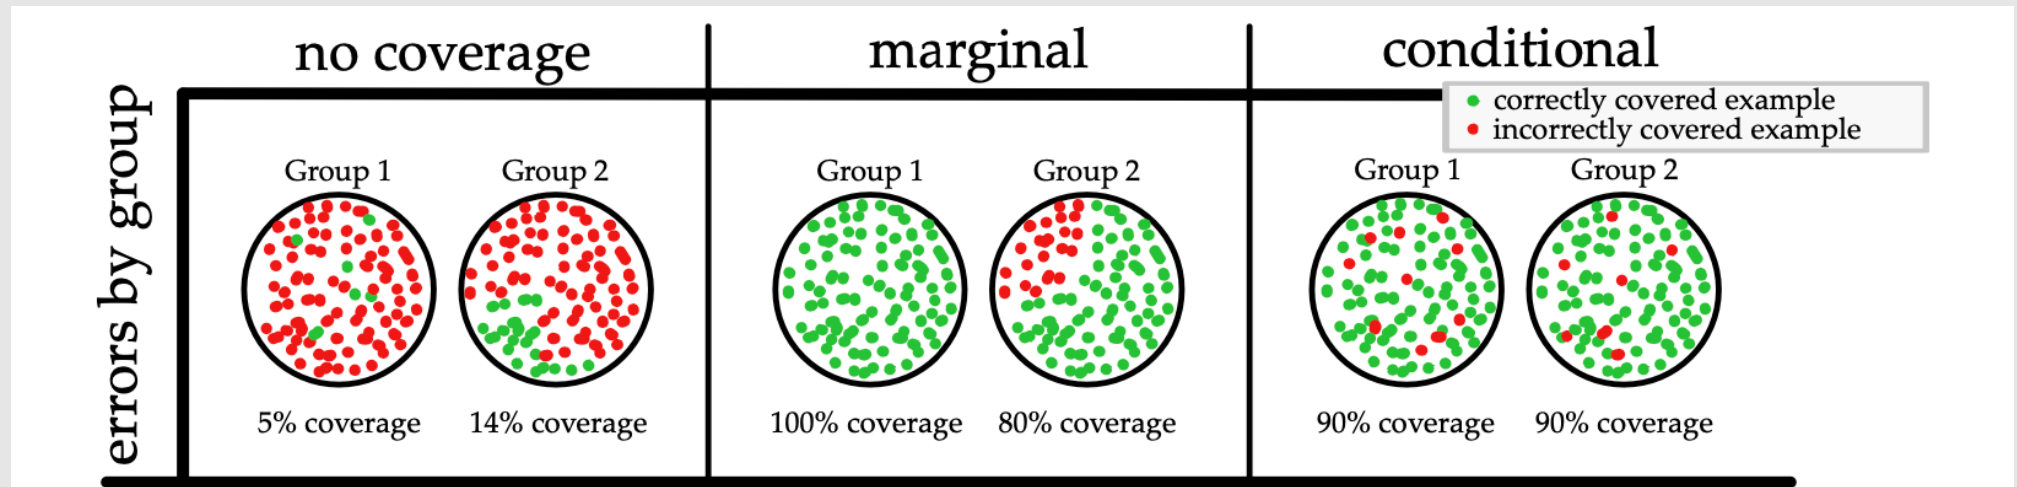


Figure from Angelopolous and Bates 2021

Devil is in the detail

- Marginal coverage guarantee **doesn't make** single predictions **particularly useful**

Marginal means “on average”: given a new set of examples that is statistically equivalent to the calibration set, the predictions sets contain the true value “x”% of the time

- What we really want is **conditional coverage**: conditional on any split of the data, i.e. for every example

Unfortunately, this is relatively easy to obtain on classes, but impossible to achieve in general

- A trivial solution to enforcing “including x% of the true classes” is to include all. **Obviously useless**
- Empirically, we look for **small** (not including too many) and **adaptive** (can recognize easy vs difficult examples; proxy for conditional coverage) prediction sets \Rightarrow look at **size** and **spread in size** of sets/intervals as diagnostics

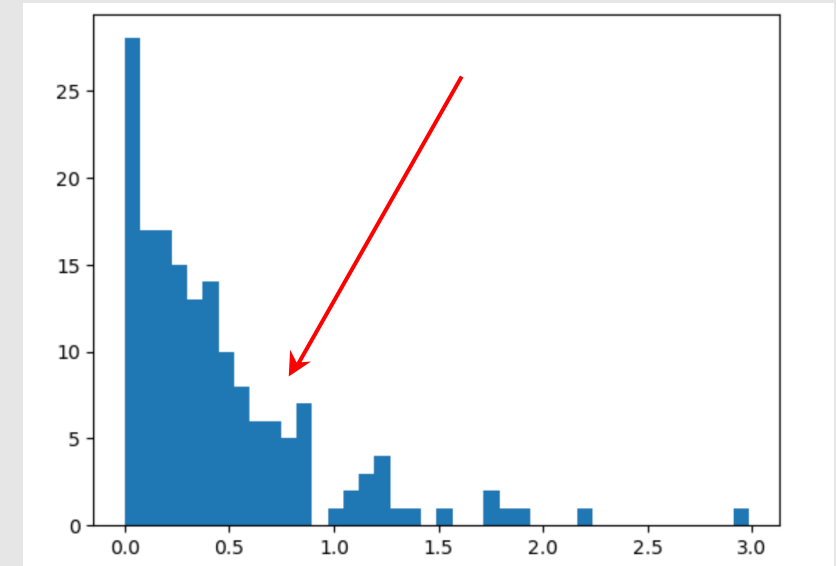
*Algorithms: **APS**, **RAPS** (Adaptive Pred Sets, Regularized Adaptive Pred Sets)*

Regression

main difference is in the **scores**

Simplest version:

- Train regression model that outputs point predictions \hat{y}
- Pick desired **coverage level α** (e.g., 80%)
- Build conformal scores on calibration set (low = good) – what could work?
- **e.g., absolute residuals $|y - \hat{y}|$**
- Find quantile of residuals **q** corresponding to desired coverage, times finite sample correction $(n+1)/n \Rightarrow \hat{q}$
- Build intervals by adding this number to/subtracting from point prediction: **$[\hat{y} - \hat{q}, \hat{y} + \hat{q}]$**
- Marginal coverage still holds!



Let's say that we are interested in 80% coverage; $\alpha = 0.2$. We find the corresponding 80% quantile of the absolute residuals, with the usual finite sample size correction:

```
alpha = 0.2
```

```
qhat = np.quantile(np.abs(y_calib_pred - y_calib), np.ceil((n+1)*(1-alpha))/n)
```

```
qhat
```

```
0.7055971501774867
```

We can now generate intervals by adding this quantity on either side of the point predictions:

```
intervals = np.hstack([(y_calib_pred - qhat).reshape(-1,1), (y_calib_pred + qhat).reshape(-1,1)])
```

To check calibration, we ask how often the true value is found in the intervals:

```
#whether true value is in interval
```

```
inint = np.array([intervals[i][0] < y_calib[i] < intervals[i][1] for i in range(len(y_calib))], dtype = int)
```

```
inint.mean() #Success! (Coverage is as expected)
```

```
0.806060606060606
```

We should also check that the coverage holds on the validation set:

```
y_val_pred = model.predict(X_val)
intervals_val = np.hstack([(y_val_pred - qhat).reshape(-1,1), (y_val_pred + qhat).reshape(-1,1)])
inint_val = np.array([intervals_val[i][0] < y_val[i] < intervals_val[i][1] for i in range(len(y_val))], dtype = int)
inint_val = np.array([intervals_val[i][0] < y_val[i] < intervals_val[i][1] for i in range(len(y_val))], dtype = int)
inint_val.mean() #Not bad! Should also marginalize over effect of split of data etc.
```

```
0.8363636363636363
```

Conformalized quantile regression

- Start from model that **outputs quantiles**
(e.g., GradientBoostingRegressor w quantile loss)
- Pick desired coverage (e.g., 80%)
- Define up/low quantiles, e.g. $\widehat{GBR}q_{\text{low}} = 0.1$; $\widehat{GBR}q_{\text{up}} = 0.9$
- On calibration set, compute **conformal scores**

$$s(x, y) = \max(\widehat{GBR}q_{\text{low}} - y, y - \widehat{GBR}q_{\text{up}})$$

- Median of s indicates typical correction needed by the intervals
(negative = intervals were too wide; positive = too narrow)
- Find quantile q of scores s , with usual correction $\Rightarrow \hat{q}$
- \hat{q} is added to both sides of the intervals **for each example**
- Adaptivity is achieved only through the original quantiles

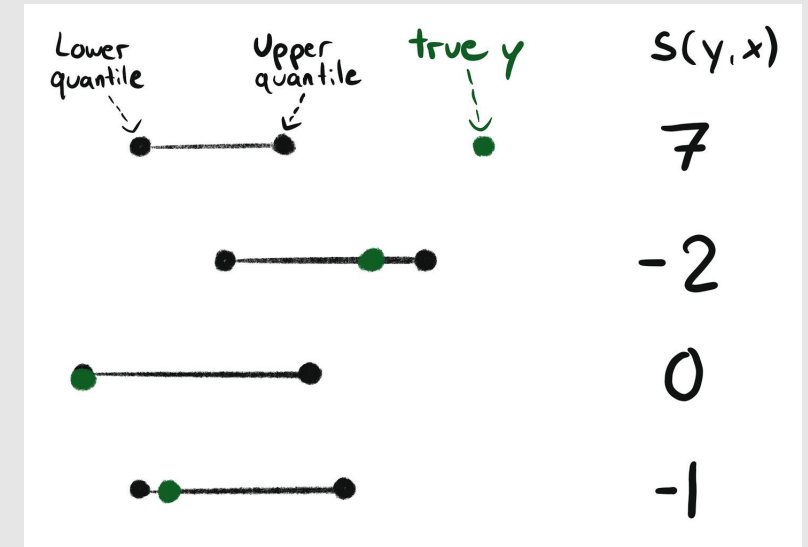


Figure by Christoph Molnar

References

- For the code relative to examples here:
<https://github.com/vacquaviva/IntroConformalPredictions>
- A paper comparing different UQ methods (including CP) for SED fitting:
<https://www.mdpi.com/2674-0346/3/1/2>
- Angelopoulos, Anastasios N., and Stephen Bates. “Conformal prediction: A gentle introduction.” Foundations and Trends® in Machine Learning 16.4 (2023): 494-591.
<https://github.com/aangelopoulos/conformal-prediction>
- Christoph Molnar has a series of (free) great blog posts and a book:
<https://mindfulmodeler.substack.com/p/week-1-getting-started-with-conformal>
<https://christophmolnar.com/books/conformal-prediction/>
- Conversation starter:
<https://statmodeling.stat.columbia.edu/2024/02/20/when-do-we-expect-conformal-prediction-sets-to-be-helpful/>
- Python Implementation:
Model Agnostic Prediction Interval Estimator: <https://mapie.readthedocs.io/en/latest/>

Summary: Uncertainty quantification methods

- Calibrating the uncertainties is always tricky. You can only do it if you have all sources of uncertainty under control. No free lunch.
- It is also necessary. Just like lunch.
- Conformal Predictions is a very active area of development (related: conformal risk control)

Methods that are proxies for conditional coverage are the most useful in practice

Question to ask: What is useful? What do you know about the data? Is it important to be conservative? Is adaptivity very important?

- Bayesian methods vs Conformal Predictions:

Do a different thing (produce calibrated sets/intervals vs produce posteriors)

- *Bayesian methods require assumptions about data distribution + priors*
- *Conformal Prediction Sets are always not always useful, even with guarantees*

- ¿Por qué no los dos?